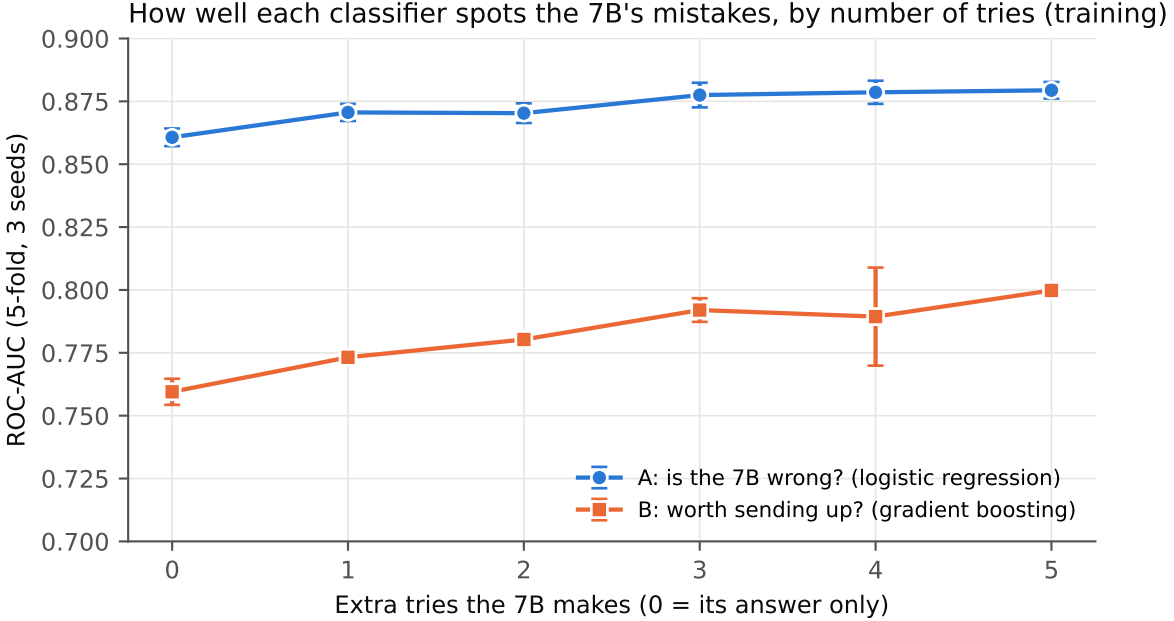
How well each classifier spots the 7B's mistakes, by number of tries (training)
ROC-AUC (5-fold, 3 seeds)
0.900
0.875
0.850
0.825
0.800
0.775
0.750
0.725
0.700
0
1
2
3
4
5
Extra tries the 7B makes (0 = its answer only)
A: is the 7B wrong? (logistic regression)
B: worth sending up? (gradient boosting)